

HANDS-ON TUTORIAL: OPEN VSWITCH* + DPDK* IN NFV USE CASE

Irene Liew

Intel Corporation

November 7th, 2016

NOTE

These slides were originally presented as part of a hands-on lab at the IEEE NFV/SDN conference in November 2016. They have been modified to make them more relevant to an audience that does not have access to the resources that were available at the time.

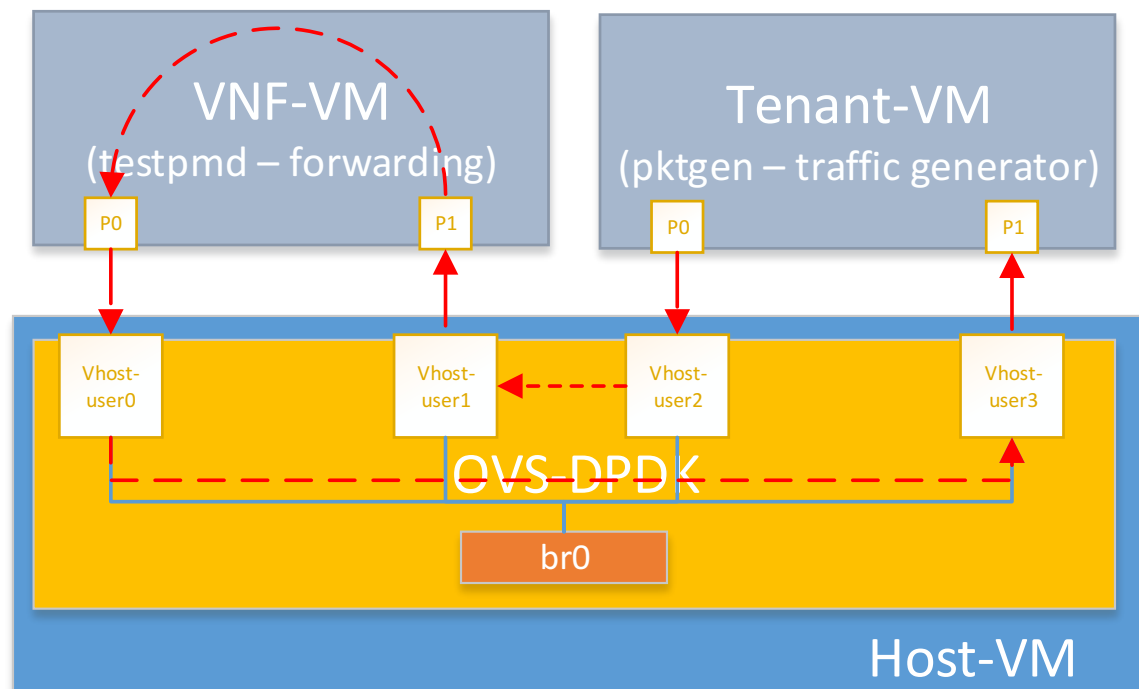


THE NEW CENTER OF POSSIBILITY

GOAL

- **Provide hands-on in DPDK* application usage in the NFV use case**
 - **Focus on the functional setup and usage in NFV use case**
 - **Setup and hands-on exercise are not optimized for performance**

OPEN VSWITCH* (OVS) + DPDK* VNF SETUP



Legend:

- ← - - - - Traffic flow
- OVS bridge connection

- Host-VM is the host machine that runs the OVS-DPDK and hosts VNF-VM and Tenant-VM:
 - a. VNF-VM: VM that would run DPDK testpmd application – packet forwarding
 - b. Tenant-VM: VM that would run pktgen application as a traffic generator
- Tenant-VM generates packets to VNF-VM, which would then be forwarded by VNF-VM back to the Tenant-VM.
- Tenant-VM shows throughput in packets/second and Mbps.
- Host-VM creates OVS bridge (br0) and 4 x vhost-user ports with each 2 ports attached to VNF-VM and Tenant-VM
- By default the 1 PMD core thread is created in OVS+DPDK. All 4 vhost-user ports are sharing a core.

LAUNCH VNF-VM & TENANT-VM

1. Login to Host-VM using this credential:
User: user
Password: password
2. After login, go to /home/user/training/dpdk-lab directory:
cd /home/user/training/dpdk-lab
ls
3. Run the following script files in order:
 - a. # sudo sh 01_start_ovs.sh //start OVS-DPDK
 - b. # sudo sh 02_createports_ovs.sh //create ports in OVS-DPDK
 - c. # sudo sh 03_addroutes_vm-vm.sh //define flows
 - d. # sudo sh 04_start_VNF-VM.sh //launch VNF-VM
 - e. # sudo sh 05_start_TenantVM.sh //launch Tenant-VM

STARTING THE EXERCISE

Requirement: putty session to login to the setup consoles.

1. Login to VNF-VM and Tenant-VM using this credential:

User: user

Password: password

2. After login to the VMs, do the following: `# sudo su -`
3. Tenant-VM shows the Pktgen application is running.
4. Testpmd application is located in VNF-VM. Go to `/root/dpdk-16.07/app/testpmd` directory:

```
# cd /root/dpdk-16.07/app/testpmd
# ls
```

```
[root@vnf-vm test-pmd]# ls
build    flowgen.c  parameters.c testpmd
cmdline.c flowgen.o  iofwd.o     parameters.o testpmd.c
cmdline.o icmpecho.c macfwd.c    _postbuild testpmd.h
config.c  icmpecho.o macfwd.o    _postinstall testpmd.map
config.o  ieee1588fwd.c macswap.c  _preinstall testpmd.o
csumonly.c _install  macswap.o  rxonly.c    txonly.c
csumonly.o iofwd.c   Makefile   rxonly.o    txonly.o
```

SIMPLE FORWARDING: IOFWD.C

- The file we will be editing is iofwd.c in function static void pkt_burst_io_forward(struct fwd_stream *fs)
vi iofwd.c
- This function forwards packets in I/O mode by forwarding packets “as-is”.
- It does not access to packets data.
- Coding instructions:
 - Implement the code for receiving packets and transmitting packets

Note:

- Receive and transmit in a burst of packets
- Use rte_eth_rx_burst and rte_eth_tx_burst function calls to implement the received and transmit
- rte_eth_* functions are defined in:
/root/dpdk-16.07/lib/librte_ether/rte_ethdev.h
- rte_mbuf structure is defined in:
/root/dpdk-16.07/lib/librte_mbuf/rte_mbuf.h

```
static void
pkt_burst_io_forward(struct fwd_stream *fs)
{
    80     struct rte_mbuf *pkts_burst[MAX_PKT_BURST];
    81     uint16_t nb_rx;
    82     uint16_t nb_tx;
    83     uint32_t retry;
    ...

    95     /*
    96      * Receive a burst of packets and forward them.
    97      */
    98     nb_rx = rte_eth_rx_burst(fs->rx_port, fs->rx_queue,
    99                             pkts_burst, nb_pkt_per_burst);
    100
    101     if (unlikely(nb_rx == 0))
    102         return;
    103     fs->rx_packets += nb_rx;
    ...

    107     /*
    108      * Transmit a burst of packets
    109      */
    110     nb_tx = rte_eth_tx_burst(fs->tx_port, fs->tx_queue,
    111                             pkts_burst, nb_rx);
    112     /*
    113      * Retry if necessary
    114      */
    if (unlikely(nb_tx < nb_rx) && fs->retry_enabled) {
        retry = 0;
        while (nb_tx < nb_rx && retry++ < burst_tx_retry_num) {
            rte_delay_us(burst_tx_delay_time);
            nb_tx += rte_eth_tx_burst(fs->tx_port, fs->tx_queue,
                                     &pkts_burst[nb_tx], nb_rx - nb_tx);
        }
    }
    fs->tx_packets += nb_tx;

    if (unlikely(nb_tx < nb_rx)) {
        fs->fwd_dropped += (nb_rx - nb_tx);
        do {
            rte_pktmbuf_free(pkts_burst[nb_tx]);
        } while (++nb_tx < nb_rx);
    }
}
```

RTE_ETH_RX_BURST

- Retrieve a burst of input packets from a receive queue of an Ethernet device. The retrieved packets are stored in **rte_mbuf** structures whose pointers are supplied in the **rx_pkts** array.

Function:

```
rte_eth_rx_burst(uint8_t port_id, uint16_t queue_id,  
                 struct rte_mbuf **rx_pkts, const uint16_t nb_pkts)
```

@param port_id: The port identifier of the Ethernet device.

@param queue_id: The index of the receive queue from which to retrieve input packets. The value must be in the range [0, nb_rx_queue - 1] previously supplied to `rte_eth_dev_configure()`. Default nb_rx_queue is 1

@param rx_pkts: The address of an array of pointers to **rte_mbuf** structures that must be large enough to store **nb_pkts** pointers in it.

@param nb_pkts: The maximum number of packets to retrieve. User defined value for packet burst size. Default value is 32.

nb_rx = the number of packets actually retrieved, which is the number of **rte_mbuf** data structures effectively supplied into the **rx_pkts** array.

```
static void  
pkt_burst_io_forward(struct fwd_stream *fs)  
{  
    80     struct rte_mbuf *pkts_burst[MAX_PKT_BURST];  
    81     uint16_t nb_rx;  
    82     uint16_t nb_tx;  
    83     uint32_t retry;  
    ...  
  
    95     /*  
    96      * Receive a burst of packets and forward them.  
    97      */  
    98     nb_rx = rte_eth_rx_burst(fs->rx_port, fs->rx_queue,  
    99                             pkts_burst, nb_pkt_per_burst);  
    100  
    101     if (unlikely(nb_rx == 0))  
    102         return;  
    103     fs->rx_packets += nb_rx;  
    ...  
}
```


RTE_ETH_TX_BURST

- Send a burst of output packets on a transmit queue of an Ethernet device.

Function:

```
rte_eth_tx_burst(uint8_t port_id, uint16_t queue_id,  
                 struct rte_mbuf **tx_pkts, uint16_t nb_pkts)
```

@param port_id: The port identifier of the Ethernet device.

@param queue_id: The index of the transmit queue through which output packets must be sent. The value must be in the range [0, nb_tx_queue - 1] previously supplied to rte_eth_dev_configure(). Default nb_tx_queue is 1.

@param tx_pkts: The address of an array of *nb_pkts* pointers to *rte_mbuf* structures which contain the output packets.

@param nb_pkts: The maximum number of packets to transmit.

nb_tx = returns the number of packets it actually sent.

```
static void  
pkt_burst_io_forward(struct fwd_stream *fs)  
{  
    80     struct rte_mbuf *pkts_burst[MAX_PKT_BURST];  
    81     uint16_t nb_rx;  
    82     uint16_t nb_tx;  
    83     uint32_t retry;  
  
    ...  
  
    ...  
    107     /*  
    108     * Transmit a burst of packets  
    109     */  
    110     nb_tx = rte_eth_tx_burst(fs->tx_port, fs->tx_queue,  
    111                             pkts_burst, nb_rx);  
    112     /*  
    113     * Retry if necessary  
    114     */  
    if (unlikely(nb_tx < nb_rx) && fs->retry_enabled) {  
        retry = 0;  
        while (nb_tx < nb_rx && retry++ < burst_tx_retry_num) {  
            rte_delay_us(burst_tx_delay_time);  
            nb_tx += rte_eth_tx_burst(fs->tx_port, fs->tx_queue,  
                                    &pkts_burst[nb_tx], nb_rx - nb_tx);  
        }  
        fs->tx_packets += nb_tx;  
  
        if (unlikely(nb_tx < nb_rx)) {  
            fs->fwd_dropped += (nb_rx - nb_tx);  
            do {  
                rte_pktmbuf_free(pkts_burst[nb_tx]);  
            } while (++nb_tx < nb_rx);  
        }  
    }
```

IOFWD.C

/root/dpdk-16.07/app/testpmd/testpmd.h:

```
/**
 * The data structure associated with a forwarding stream between a receive
 * port/queue and a transmit port/queue.
 */
struct fwd_stream {
    /* "read-only" data */
    portid_t rx_port; /**< port to poll for received packets */
    queueid_t rx_queue; /**< RX queue to poll on "rx_port" */
    portid_t tx_port; /**< forwarding port of received packets */
    queueid_t tx_queue; /**< TX queue to send forwarded packets */
    streamid_t peer_addr; /**< index of peer ethernet address of packets */

    unsigned int retry_enabled;

    /* "read-write" results */
    unsigned int rx_packets; /**< received packets */
    unsigned int tx_packets; /**< received packets transmitted */
    unsigned int fwd_dropped; /**< received packets not forwarded */
    unsigned int rx_bad_ip_csum; /**< received packets has bad ip checksum */
    unsigned int rx_bad_l4_csum; /**< received packets has bad l4 checksum */
#ifdef RTE_TEST_PMD_RECORD_CORE_CYCLES
    uint64_t core_cycles; /**< used for RX and TX processing */
#endif
#ifdef RTE_TEST_PMD_RECORD_BURST_STATS
    struct pkt_burst_stats rx_burst_stats;
    struct pkt_burst_stats tx_burst_stats;
#endif
};
```

```
static void
pkt_burst_io_forward(struct fwd_stream *fs)
{
    80     struct rte_mbuf *pkts_burst[MAX_PKT_BURST];
    81     uint16_t nb_rx;
    82     uint16_t nb_tx;
    83     uint32_t retry;
    ...

    95     /*
    96     * Receive a burst of packets and forward them.
    97     */
    98     nb_rx = rte_eth_rx_burst(fs->rx_port, fs->rx_queue,
    99                             pkts_burst, nb_pkt_per_burst);
    100
    101     if (unlikely(nb_rx == 0))
    102         return;
    103     fs->rx_packets += nb_rx;
    ...

    107     /*
    108     * Transmit a burst of packets
    109     */
    110     nb_tx = rte_eth_tx_burst(fs->tx_port, fs->tx_queue,
    111                             pkts_burst, nb_rx);
    112     /*
    113     * Retry if necessary
    114     */
    if (unlikely(nb_tx < nb_rx) && fs->retry_enabled) {
        retry = 0;
        while (nb_tx < nb_rx && retry++ < burst_tx_retry_num) {
            rte_delay_us(burst_tx_delay_time);
            nb_tx += rte_eth_tx_burst(fs->tx_port, fs->tx_queue,
                                     &pkts_burst[nb_tx], nb_rx - nb_tx);
        }
    }
    fs->tx_packets += nb_tx;

    if (unlikely(nb_tx < nb_rx)) {
        fs->fwd_dropped += (nb_rx - nb_tx);
        do {
            rte_pktmbuf_free(pkts_burst[nb_tx]);
        } while (++nb_tx < nb_rx);
    }
}
```

TIPS IN RUNNING TESTPMD IN VNF-VM

1. Save the iofwd.c file:
Press "Esc" key
#:wq!
2. Go to /root/dpdk-16.07/app/testpmd and compile testpmd application:
cd /root/dpdk-16.07/app/testpmd
#export RTE_SDK=/root/dpdk-16.07
#export RTE_TARGET=x86_64-native-linuxapp-gcc
#make
2. Start testpmd application by running:
#sh /root/run_testpmd.sh
3. Once testpmd application is running, do the following:
testpmd> set fwd io
testpmd> start

```
PMD: bnxt rte_pmd_init() called for (null)
EAL: PCI device 0000:00:03.0 on NUMA socket -1
EAL: probe driver: 1af4:1000 rte_virtio_pmd
EAL: PCI device 0000:00:04.0 on NUMA socket -1
EAL: probe driver: 1af4:1000 rte_virtio_pmd
EAL: PCI device 0000:00:05.0 on NUMA socket -1
EAL: probe driver: 1af4:1000 rte_virtio_pmd
Interactive-mode selected
USER1: create a new mbuf pool <mbuf_pool_socket_0>: n=155456, size=2176, socket=
0
Configuring Port 0 (socket 0)
Port 0: 00:00:00:00:00:01
Configuring Port 1 (socket 0)
Port 1: 00:00:00:00:00:02
Checking link statuses...
Port 0 Link Up - speed 10000 Mbps - full-duplex
Port 1 Link Up - speed 10000 Mbps - full-duplex
Done
testpmd> set fwd io
Set io packet forwarding mode
testpmd> start
io packet forwarding - ports=2 - cores=1 - streams=2 - NUMA support disabled, MP
over anonymous pages disabled
Logical Core 2 (socket 0) forwards packets on 2 streams:
RX P=0/Q=0 (socket 0) -> TX P=1/Q=0 (socket 0) peer=02:00:00:00:00:01
RX P=1/Q=0 (socket 0) -> TX P=0/Q=0 (socket 0) peer=02:00:00:00:00:00

io packet forwarding - CRC stripping disabled - packets/burst=64
nb forwarding cores=1 - nb forwarding ports=2
RX queues=1 - RX desc=2048 - RX free threshold=0
RX threshold registers: pthresh=0 hthresh=0 wthresh=0
TX queues=1 - TX desc=2048 - TX free threshold=0
TX threshold registers: pthresh=0 hthresh=0 wthresh=0
TX RS bit threshold=0 - TXQ flags=0xf00
testpmd> show port stats all
```

RUNNING PKTGEN IN TENANT-VM

Let's start generating 100 packets from Port 0:

1. Pktgen > set 0 count 100
2. To start traffic on Port0:
Pktgen > start 0

To start pktgen to generate 10% line rate from Port 0:

1. Pktgen > set 0 rate 10
2. To start traffic on Port0:
Pktgen > start 0

```
- Ports 0-1 of 2 <Main Page> Copyright (c) <2010-2016>, Intel Corporation
Flags:Port : P-----:0 P-----:1
Link State : <UP-10000-FD> <UP-10000-FD> ----TotalRate----
Pkts/s Max/Rx : 0/0 100/0 100/0
Max/Tx : 100/0 0/0 100/0
Mbits/s Rx/Tx : 0/0 0/0 0/0
Broadcast : 0 0
Multicast : 0 0
64 Bytes : 0 100
65-127 : 0 0
128-255 : 0 0
256-511 : 0 0
512-1023 : 0 0
1024-1518 : 0 0
Runts/Jumbos : 0/0 0/0
Errors Rx/Tx : 0/0 0/0
Total Rx Pkts : 0 100
Tx Pkts : 100 0
Rx MBs : 0 0
Tx MBs : 0 0
ARP/ICMP Pkts : 0/0 0/0
Pattern Type : abcd... abcd...
Tx Count/% Rate : 100 / 100% Forever / 100%
PktSize/Tx Burst : 64 / 32 64 / 32
Src/Dest Port : 1234 / 5678 1234 / 5678
Pkt Type:VLAN ID : IPv4 / TCP:0001 IPv4 / TCP:0001
Dst IP Address : 192.168.1.1 192.168.0.1
Src IP Address : 192.168.0.1/24 192.168.1.1/24
Dst MAC Address : 00:00:00:00:00:04 00:00:00:00:00:03
Src MAC Address : 00:00:00:00:00:03 00:00:00:00:00:04
VendID/PCI Addr : 1af4:1000/00:04.0 1af4:1000/00:05.0

-- Pktgen Ver: 3.0.14 (DPDK 16.07.0) Powered by Intel® DPDK -----

Pktgen > set 0 count 100
Pktgen> start 0
Pktgen>
```

HEX DUMP OF PACKETS IN PKTGEN

- Default packet pattern type is alphabet string (abcd...)
- To view the hex dump of the first packet to be sent on Port 0, type:

Pktgen > pdump 0

```
00000000: 00 00 00 00 00 04 00 00 00 00 00 03 08 00 45 00 | .....E.  
00000010: 00 2E 18 37 00 00 04 06 1C 41 C0 A8 00 01 C0 A8 | ...7....A.....  
00000020: 01 01 04 D2 16 2E 12 34 56 78 12 34 56 90 50 10 | .....4Vx.4V.P.  
00000030: 20 00 FF E6 00 00 77 78 79 7A 30 31 | | | | .....wxyz01
```

OTHER TIPS IN PKTGEN RUNTIME COMMAND LINE

- To clear the statistics for all ports, type:
`Pktgen > clr`
- To reset the configuration for all ports, type:
`Pktgen > rst`
- To set DEST MAC address on Port 0, type:
`Pktgen > set mac 0 00:00:00:00:00:02`
- Other Pktgen options:
`Pktgen> help`

MONITOR TRAFFIC IN VNF-VM

- To monitor the Tx and Rx traffic in the testpmd application:

testpmd> show port stats all

```
testpmd> show port stats all

##### NIC statistics for port 0 #####
RX-packets: 0      RX-missed: 0      RX-bytes: 0
RX-errors: 0
RX-nobuf: 0
TX-packets: 100    TX-errors: 0      TX-bytes: 6000

Throughput (since last show)
Rx-pps: 0
Tx-pps: 368695647
#####

##### NIC statistics for port 1 #####
RX-packets: 100    RX-missed: 0      RX-bytes: 6000
RX-errors: 0
RX-nobuf: 0
TX-packets: 0      TX-errors: 0      TX-bytes: 0

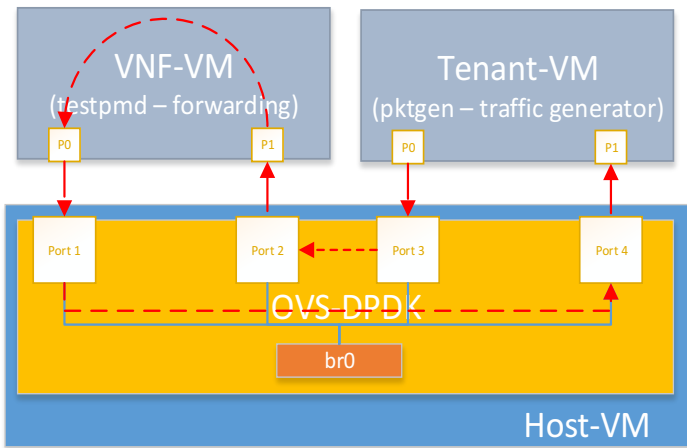
Throughput (since last show)
Rx-pps: 368695605
Tx-pps: 0
#####
```

Note: 100 packets are received at Port 1 and transmit out from Port 0

- To clear the ports statistics in the testpmd application:

testpmd> clear port stats all

WORKING VNF FORWARDING APPLICATION: FWD IO



Legend:

- ← - - - Traffic flow
- OVS bridge connection

Tenant-VM

```
Flags:Port : P-----:0 P-----:1
Link State : P-----:0 P-----:1
<UP-10000-FD> <UP-10000-FD>
10944/8224 11648/7744
726336/629632 710464/641216
5/423 5/430
64 Bytes : 0 0
65-127 : 6328576 6353512
128-255 : 0 0
256-511 : 0 0
512-1023 : 0 0
1024-1518 : 0 0
Runt/Jumbos : 0 0
Errors Rx/Tx : 0/0 0/0
Total Rx Pkts : 0/0 0/0
Tx Pkts : 6325248 6350696
Rx MBs : 505662272 507402720
Tx MBs : 4250 4267
ARP/ICMP Pkts : 339805 340974
Pattern Type : abcd... abcd...
Tx Count/% Rate : Forever / 100% Forever / 100%
PktSize/Tx Burst : 64 / 32 64 / 32
Src/Dst Port : 1234 / 5678 1234 / 5678
Pkt Type:VLAN ID : IPv4 / TCP:0001 IPv4 / TCP:0001
Dst IP Address : 192.168.1.1 192.168.0.1
Src IP Address : 192.168.0.1/24 192.168.1.1/24
Dst MAC Address : 00:00:00:00:00:04 00:00:00:00:00:03
Src MAC Address : 00:00:00:00:00:03 00:00:00:00:00:04
VendID/PCI Addr : 1af4:1000/00:04.0 1af4:1000/00:05.0

-- Pktgen Ver: 3.0.14 (DPDK 16.07.0) Powered by Intel® DPDK -----
Pktgen > start all
Pktgen>
```

VNF-VM

```
##### NIC statistics for port 1 #####
RX-packets: 211507492 RX-missed: 0 RX-bytes: 12690451212
RX-errors: 0
RX-nombuf: 0
TX-packets: 212011174 TX-errors: 0 TX-bytes: 12720671850

Throughput (since last show)
Rx-pps: 354098
Tx-pps: 354942
#####
testpmd>
testpmd> show port stats all

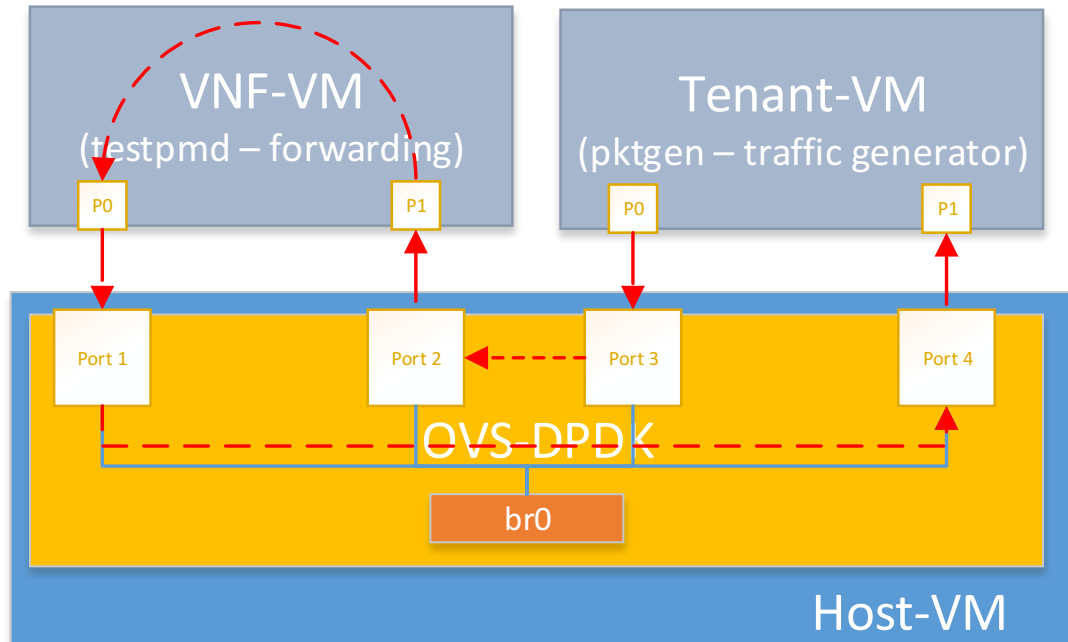
##### NIC statistics for port 0 #####
RX-packets: 331695972 RX-missed: 0 RX-bytes: 19901759730
RX-errors: 0
RX-nombuf: 0
TX-packets: 330449052 TX-errors: 0 TX-bytes: 19826944812

Throughput (since last show)
Rx-pps: 620370
Tx-pps: 617297
#####
##### NIC statistics for port 1 #####
RX-packets: 330569764 RX-missed: 0 RX-bytes: 19834187532
RX-errors: 0
RX-nombuf: 0
TX-packets: 331570086 TX-errors: 0 TX-bytes: 19894206570

Throughput (since last show)
Rx-pps: 617547
Tx-pps: 620123
#####
testpmd>
```

- Pktgen in Tenant-VM generates packets from Port 0 and transmits them to VNF-VM via vhost-user Port 3 -> Port 2 running in OVS-DPDK (Host-VM).
- A working testpmd application in VNF-VM receives the packets at P1 and transmits them out at P0 to Tenant-VM via vhost-user Port 1 -> Port 4
- Observed transmit (Tx) packets and receive (Rx) packets on the Tenant-VM and VNF-VM ports.

PORT NUMBERING IN OVS AND VM



- Port numbering in the OVS Host-VM is shown via command line:

`#/home/user/ovs/utilities/ovs-ofctl show br0`

- Output screen:

```
1(vhost-user0): addr:00:00:00:00:00:00
  config:  PORT_DOWN
  state:   LINK_DOWN
  speed: 0 Mbps now, 0 Mbps max
2(vhost-user1): addr:00:00:00:00:00:00
  config:  PORT_DOWN
  state:   LINK_DOWN
  speed: 0 Mbps now, 0 Mbps max
3(vhost-user2): addr:00:00:00:00:00:00
  config:  PORT_DOWN
  state:   LINK_DOWN
  speed: 0 Mbps now, 0 Mbps max
4(vhost-user3): addr:00:00:00:00:00:00
  config:  PORT_DOWN
  state:   LINK_DOWN
  speed: 0 Mbps now, 0 Mbps max
```

- In each VNF-VM and Tenant-VM view, the ports are shown as Port 0 and Port 1

PACKET MONITORING IN OVS-DPDK: HOST-VM

1. Check the vhost-user ports' stats (e.g. RX, Tx, drop, error, etc) in Host-VM:

- `#!/home/user/ovs/utilities/ovs-ofctl dump-ports br0`
- Run the following script to get the real-time stats:
`# sh /home/user/training/dump-ports.sh`

Every 1.0s: `/root/ovs/utilities/ovs-ofctl dump-ports br0`

```
OFPST_PORT reply (xid=0x2): 5 ports
port 4: rx pkts=38467002, bytes=2308023618, drop=0, errs=0, frame=?, over=?, crc=?
      tx pkts=526754, bytes=31630962, drop=37265166, errs=?, coll=?
port 2: rx pkts=37775940, bytes=2266589652, drop=0, errs=0, frame=?, over=?, crc=?
      tx pkts=37801705, bytes=2268105744, drop=251041, errs=?, coll=?
port 1: rx pkts=37792073, bytes=2267553120, drop=0, errs=0, frame=?, over=?, crc=?
      tx pkts=37785865, bytes=2267155062, drop=681120, errs=?, coll=?
port 3: rx pkts=38052763, bytes=2283169560, drop=0, errs=0, frame=?, over=?, crc=?
      tx pkts=520202, bytes=31242354, drop=37255585, errs=?, coll=?
port LOCAL: rx pkts=0, bytes=0, drop=0, errs=0, frame=0, over=0, crc=0
           tx pkts=0, bytes=0, drop=0, errs=0, coll=0
```

2. Check the traffic flows among the vhost-user ports in Host-VM:

- `#!/home/user/ovs/utilities/ovs-ofctl dump-flows br0`
- Run the following script to get the real-time stats:
`# sh /home/user/training/dump-flows.sh`

Every 1.0s: `/root/ovs/utilities/ovs-ofctl dump-flows br0`

```
NXST_FLOW reply (xid=0x4):
cookie=0x0, duration=2976.679s, table=0, n_packets=0, n_bytes=0, idle_age=0, ip,in_port=2 actions=output:3
cookie=0x0, duration=2976.677s, table=0, n_packets=158970348, n_bytes=9538224264, idle_age=0, ip,in_port=3 actions=output:2
cookie=0x0, duration=2976.676s, table=0, n_packets=156882706, n_bytes=9412988304, idle_age=0, ip,in_port=1 actions=output:4
cookie=0x0, duration=2976.674s, table=0, n_packets=0, n_bytes=0, idle_age=0, ip,in_port=4 actions=output:1
```

REREFENCES

- DPDK Programmer's Guide: http://dpdk.org/doc/guides/prog_guide/index.html
- Testpmd Application User Guide: http://dpdk.readthedocs.io/en/v16.04/testpmd_app_ug/index.html
- Sample Application User Guide: http://dpdk.readthedocs.io/en/v16.04/sample_app_ug/index.html
- ONP 2.1 Performance Test Report (https://download.01.org/packet-processing/ONPS2.1/Intel_ONP_Release_2.1_Performance_Test_Report_Rev1.0.pdf)
- How to get best performance with NICs on Intel platforms with DPDK: http://dpdk.org/doc/guides-2.2/linux_gsg/nic_perf_intel_platform.html
- Open vSwitch 2.5.0 documentation and installation guide:
 - <http://openvswitch.org/support/dist-docs-2.5/>
 - <https://github.com/openvswitch/ovs/blob/master/INSTALL.DPDK-ADVANCED.md>